Quarterly Operations Report: Document Intake

This report summarizes how incoming documents were processed in the third quarter, where the pipeline lost time, and which changes are planned for the next quarter.

1. Volume

The intake service received 18,420 documents, 12 percent more than in the second quarter. Invoices made up 41 percent of the volume, contracts 23 percent and receipts 19 percent.

Peak days followed the month-end closing, when the queue grew to four times its average length.

2. Processing Time

The median time from upload to a ready document fell from 96 seconds to 71 seconds. Scanned documents remained the slowest group because every page passes through optical character recognition.

Field extraction accounted for most of the remaining time on long contracts, where the whole text is sent to the model in one request.

3. Review Rate

Operators reviewed 8.4 percent of the documents, down from 11.2 percent. Most reviews were caused by missing invoice numbers on photographed receipts and by unclear contract parties.

Reviews took a median of three minutes each, and nine out of ten ended with a correction to a single field.

4. Next Quarter

The team will run independent stages in parallel, retry throttled model calls with backoff, and measure parsing quality on a fixed corpus before every release.

A layout-aware parser will replace the plain text layer for scanned pages, so that answers can cite the page they come from and reviewers can see the region a value was read from.

The review queue will be split by document type, so that the operators who know a type best receive its documents first.